

1) Get the List of Open weight LLM from different service providers

a) Open AI:

GPT-2, GPT-OSS-120b, GPT-OSS-20b & GPT-OSS-safeguard

b) QWEN:

- Qwen3, Qwen2.5, Qwen2 series
- Qwen3.5-397B-A17B plus smaller Qwen3.5 sizes like 9B (Feb 2026)
- Qwen3.6-27B and Qwen3.6-35B-A3B (Apr 2026)
- Qwen3.8-2.4T-A95B and Qwen3.8-27B (Aug 2026)
- Qwen3.8-Flash-Next (Aug 2026)

c) DeepSeek:

- DeepSeek-V3 / V3.1 / V3.2, DeepSeek-R1, DeepSeek-V2, DeepSeek-Coder
- DeepSeek-V4-Pro and DeepSeek-V4-Flash (original April 2026 releases)
- DeepSeek-V4-Pro-0813 and DeepSeek-V4-Flash-0731 (Jul–Aug 2026 refreshes)

2) Hallucination check for the below prompt

Write a 4-line poem where every line ends with a word that strictly rhymes with the word 'Silver'. In all AI Model

a) QWEN AI:

The moonlight dances on the river,
Making the autumn breezes shiver,
Watching the leafy branches quiver,
Peace is a gentle, quiet giver.

It is given the response but not Properly .

b) CHATGBT:

A strict, standard English rhyme for “**silver**” essentially does not exist. So a 4-line poem where **every line ends in a strict perfect rhyme with “silver”** cannot be written without using obscure/nonstandard words.

If you allow **near rhyme**, I can make one using words like *shiver*, *river*, and *quiver*.

c) Gemini:

A quiet wind begins to shivering,
Along the path where shadows sliver,
The moonlight makes the river quiver,
And softly sets the leaves a-thliver.

d) Claude:

"Silver" famously has almost no strict rhymes in English. The only real one is **chilver**, an old dialect word for a young female lamb. So a strictly rhyming 4-line poem has to reuse it:

*The moon spilled light upon a chilver,
Asleep beside her mother, chilver,
The shepherd whispered, "Rest now, chilver,"
And stars kept watch above the chilver.*

If you'd rather have four different ending words, I can write one with near-rhymes like shiver, kilter, pilfer, and filter.

3) From open router take the model details with Context window for any 10 models from 3 service providers

Objective: Compare 10 LLMs from 3 providers (OpenAI, DeepSeek, Qwen) by context window and pricing.

#	Provider	Model	Context Window	Input / Output (\$ per 1M tokens)	Weights
1	OpenAI	GPT-6 Astra	1.05M	\$10 / \$50	Closed
2	OpenAI	GPT-5.6 Sol	1.05M	\$2 / \$10	Closed
3	OpenAI	GPT-5.4 Mini	400K	\$0.75 / \$4.50	Closed
4	OpenAI	gpt-oss-120b	131K	\$0.03 / \$0.17	Open
5	DeepSeek	DeepSeek V4.1 Flash	1.05M	\$0.15 / \$1.20	Open
6	DeepSeek	DeepSeek V4 Pro 0813	1.05M	\$0.99 / \$2.97	Open
7	DeepSeek	DeepSeek V3.2	164K	\$0.21 / \$0.31	Open
8	Qwen	Qwen3.8 Max (0902)	1M	\$2 / \$6	Closed (API)
9	Qwen	Qwen3.8 27B	1M	\$0.15 / \$2	Open
10	Qwen	Qwen3.6 35B A3B	262K	\$0.05 / \$0.70	Open

Reference model: gpt-oss-120b

- Official specs (OpenAI docs): 131,072 context window, 131,072 max output tokens, Jun 01, 2024 knowledge cutoff. It has 117B parameters with 5.1B active parameters and fits on a single H100 GPU. [openai/openai](#)
- Running it locally (Ollama): The download is 65GB, MXFP4 quantization, and the run command is `ollama run gpt-oss:120b`. [ollama](#)

Tokenizer note (platform.openai.com/tokenizer)

Context window is measured in tokens, not words. A common rule of thumb for English is $1 \text{ token} \approx 4 \text{ characters} \approx 0.75 \text{ words}$. At that rate, a 1M-token window holds about 750K words, and gpt-oss-120b's 131K window holds about 98K words.

Observation

Most current flagship models offer a context window of about 1M tokens. Smaller or older open-weight models like gpt-oss-120b and DeepSeek V3.2 are in the 131K–262K range but cost much less.